

Simple Web Application Deployment Project on AWS

Please open each step by clicking on ➔ and see the explanations.

Step 1: Create a Simple Web Application

I created a basic web app using Flask. The app takes two numbers as input and calculates the formula:

$$\sqrt{x^2 + y^2} + \sin(x * y)$$

It then returns the result to the web page. The app has only two files, `app.py` (Python code) and one HTML file. No CSS or images were used to keep it simple.

The app runs locally (on my laptop) and works correctly.

In command prompt:

```
Administrator: Command Prompt - python app.py
Microsoft Windows [Version 10.0.19045.6456]
(c) Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>cd C:\Users\msaam\Desktop\Docker

C:\Users\msaam\Desktop\Docker>python app.py
* Serving Flask app 'app'
* Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on http://127.0.0.1:5000
Press CTRL+C to quit
* Restarting with stat
* Debugger is active!
* Debugger PIN: 634-097-475
```

In browser:



Step 2: Containerization with Docker

I tested three methods: Docker Desktop, [Play with Docker](#), and GitHub. I chose GitHub because it connects easily to AWS ECR. (I could also build docker image inside AWS through CodeBuild)

Note:

In `app.py`, I updated this line:

```
```python
app.run(debug=True)
```
to
```python
app.run(host='0.0.0.0', port=5000, debug=True)
```

```

This change makes the app accessible through the cloud instead of only localhost and application in deployment not be stranded because of the address.

I created a public GitHub repository named **MathApp** ([Github Repo - MathApp](#)), added a `Dockerfile`, and a `requirements.txt` file listing dependencies. These files help others understand and reproduce the setup.

Step 3: Push Docker Image to AWS ECR

I created a new repository in AWS ECR called **mathapp-repo** (it can be also done with lac and bash file). Then, in GitHub Actions, I added a workflow file `docker-to-ecr.yml` to build and push the Docker image to AWS ECR.

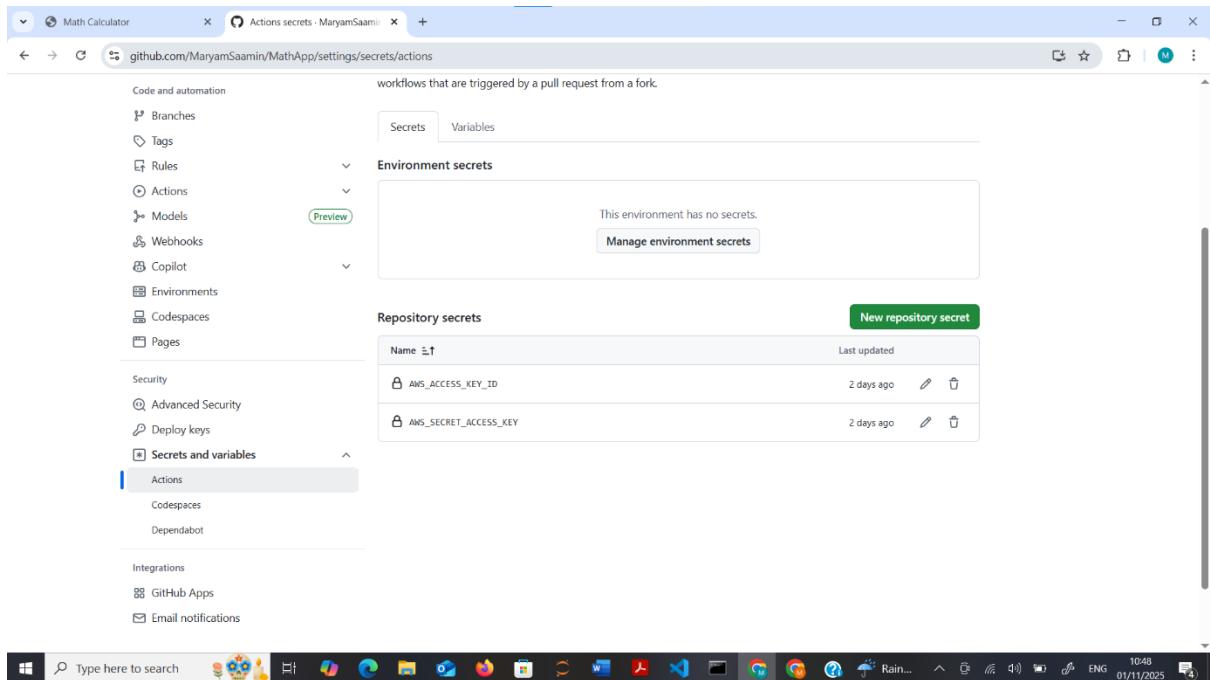
Note:

Before running the workflow, I added AWS credentials in:

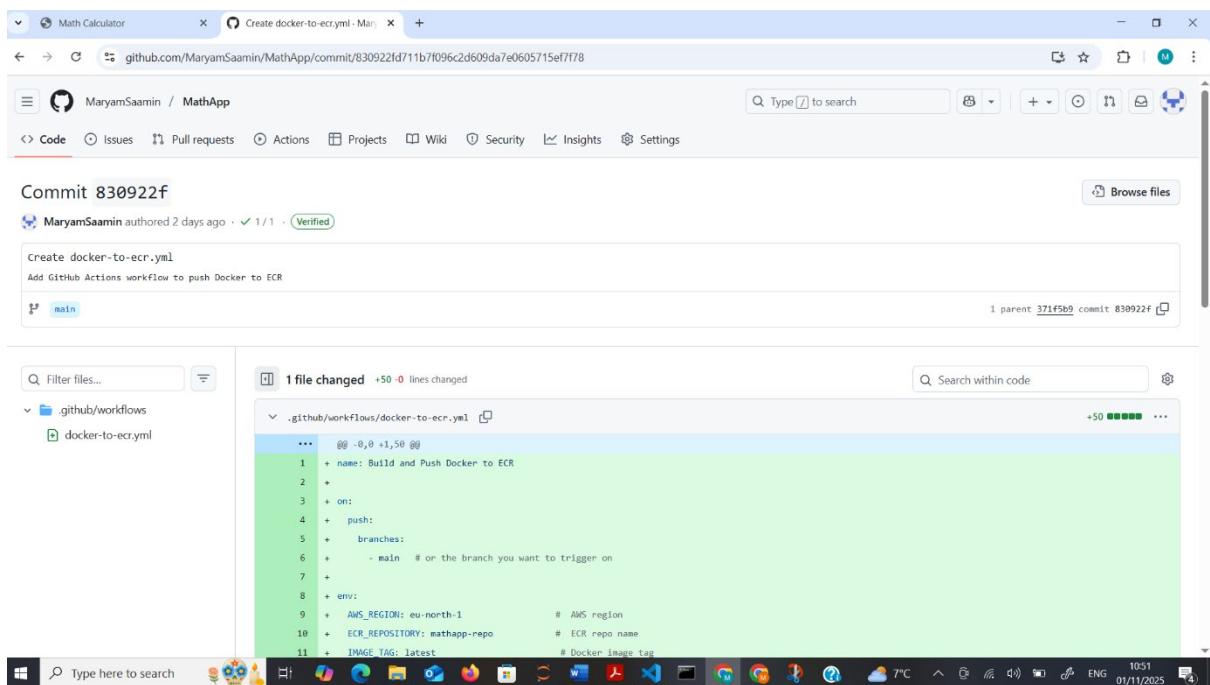
`Settings > Secrets and Variables > Actions > New Repository Secret`

(`AWS_ACCESS_KEY_ID`, `AWS_SECRET_ACCESS_KEY`), otherwise push action will be failed.

The workflow ran successfully.



The screenshot shows the GitHub Actions secrets settings page for a repository named "MathApp". The left sidebar includes sections for Code and automation, Security, Integrations, and Actions. Under Actions, the "Secrets and variables" section is selected. The main content area displays two secret types: Environment secrets and Repository secrets. There are no environment secrets listed. Under Repository secrets, there are two entries: "AWS_ACCESS_KEY_ID" and "AWS_SECRET_ACCESS_KEY", both updated 2 days ago.



The screenshot shows a GitHub commit details page for commit 830922f. The commit message is "Create docker-to-ecr.yml - Main". The commit was authored by MaryamSaamin 2 days ago and is verified. The commit details show a single file change: ".github/workflows/docker-to-ecr.yml". The code editor displays the YAML configuration for the workflow:

```
name: Build and Push Docker to ECR
on:
  push:
    branches:
      - main # or the branch you want to trigger on
  env:
    AWS_REGION: eu-north-1 # AWS region
    ECR_REPOSITORY: mathapp-repo # ECR repo name
    IMAGE_TAG: latest # Docker image tag
```

And the Docker image was uploaded to ECR.

The screenshot shows the AWS ECR console interface. On the left, there's a sidebar with navigation links for Amazon Elastic Container Registry (Private registry, Public registry, ECR public gallery, Amazon ECS, Amazon EKS, Getting started, Documentation). The main area is titled "Private repositories (2)" and lists two entries:

| Repository name | URI | Created at | Tag immutability | Encryption type |
|--|--|------------------------------------|------------------|-----------------|
| cdk-hnb659fds-container-assets-593970662859-eu-north-1 | 593970662859.dkr.ecr.eu-north-1.amazonaws.com/cdk-hnb659fds-container-assets-593970662859-eu-north-1 | 31 October 2025, 10:34:34 (UTC+02) | Immutable | AES-256 |
| mathapp-repo | 593970662859.dkr.ecr.eu-north-1.amazonaws.com/mathapp-repo | 30 October 2025, 15:17:09 (UTC+02) | Mutable | AES-256 |

At the bottom of the page, there's a CloudShell tab and a footer with links for Actions, Terms, and Cookie preferences.

Step 4: Deploy with AWS CDK (Infrastructure as Code)

Next, I deployed the web app on AWS using AWS CDK with Python.

Steps:

1. Install AWS CDK:

```
```bash
npm install -g aws-cdk
cdk --version
````
```

2. Create and initialize the CDK project:

```
```bash
mkdir flask-cdk
cd flask-cdk
cdk init app --language python
````
```

3. Set up Python environment:

```
```bash
python3 -m venv .env
source .env/bin/activate
pip install -r requirements.txt
pip install "aws-cdk-lib>=2.0.0" "constructs>=10.0.0,<11.0.0"
````
```

4. Edit the file `flask_cdk/flask_cdk_stack.py` with the following stack configuration (creating VPC, ECS cluster, and Fargate service with a load balancer).

Open my stack file

```
nano flask_cdk/flask_cdk_stack.py
```

Replace everything with this:

```
from aws_cdk import (
    Stack,
    aws_ec2 as ec2,
    aws_ecs as ecs,
    aws_ecs_patterns as ecs_patterns,
    aws_iam as iam,
    CfnOutput,
)

from constructs import Construct

class FlaskCdkStack(Stack):

    def __init__(self, scope: Construct, construct_id: str, **kwargs) -> None:
        super().__init__(scope, construct_id, **kwargs)

        # Create a new VPC
        vpc = ec2.Vpc(self, "FlaskVpc", max_azs=2)

        # Create ECS Cluster
        cluster = ecs.Cluster(self, "FlaskCluster", vpc=vpc)

        # Create ECS Task Execution Role with proper permissions
        execution_role = iam.Role(
            self, "FlaskTaskExecutionRole",
            assumed_by=iam.ServicePrincipal("ecs-tasks.amazonaws.com")
        )
        execution_role.add_managed_policy(
```

```

        iam.ManagedPolicy.from_aws_managed_policy_name(
            "service-role/AmazonECSTaskExecutionRolePolicy"
        )
    )

# Create Fargate service behind a Load Balancer
flask_service = ecs_patterns.ApplicationLoadBalancedFargateService(
    self,
    "FlaskService",
    cluster=cluster,
    cpu=256,
    memory_limit_mib=512,
    desired_count=1,
    public_load_balancer=True,
    task_image_options=ecs_patterns.ApplicationLoadBalancedTaskImageOptions(
        image=ecs.ContainerImage.from_registry(
            "593970662859.dkr.ecr.eu-north-1.amazonaws.com/mathapp-repo:latest"
        ),
        container_port=5000,
        execution_role=execution_role #<-- assign role here
    ),
)
```
Output Load Balancer DNS
CfnOutput(self, "LoadBalancerURL",
 value=f"http://{{flask_service.load_balancer.load_balancer_dns_name}}"
)

```

5. After editing, I deployed it with:

```

```bash
cdk bootstrap
cdk deploy

```

...

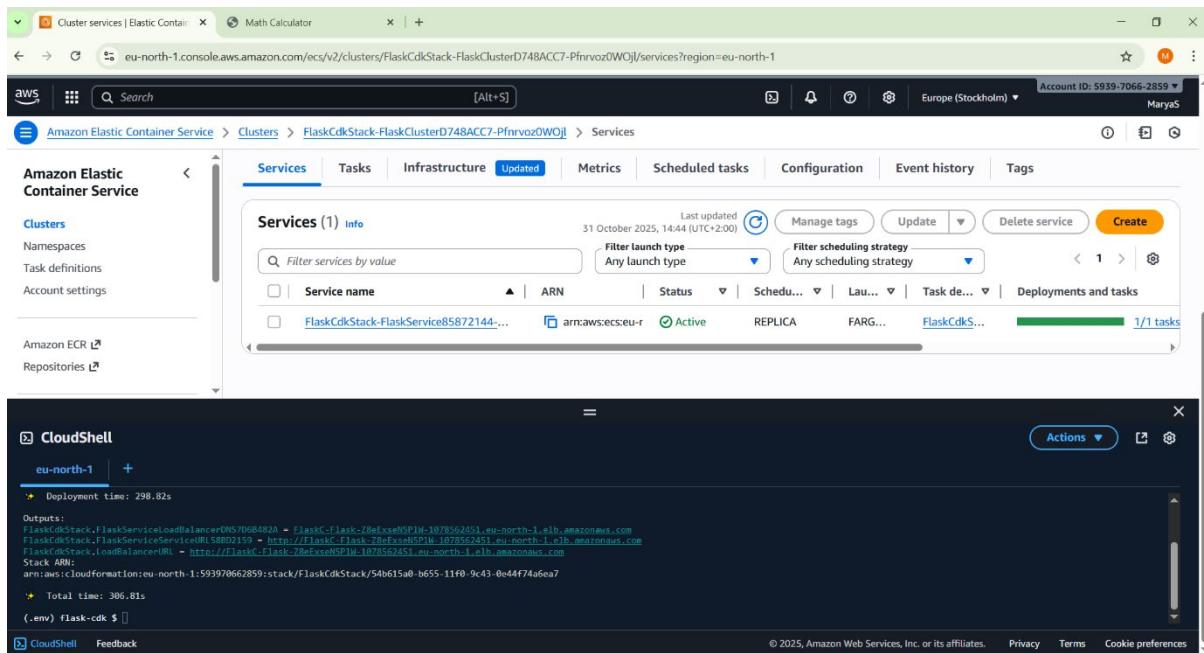
Note:

When the first deployment got stuck, I checked ECS Service events and found a missing IAM permission. After fixing the role permissions in the stack file, redeploying worked successfully (I went to ECS Service events, ECS → Clusters → Services → FlaskService → Events and I saw one task stopped and checked the failing to pull the Docker image from ECR because the task execution role is missing permissions and I added to stack file.)

The web app was live at:

AWS Console → CloudFormation → Output

<http://flaskc-flask-z8eexsen5plw-1078562451.eu-north-1.elb.amazonaws.com/>





All the steps of deployment can be executed using the **deploy.sh** script. I initially ran them step by step, and then combined them into a single batch file.

Step 5: Clean Up Resources

To delete all AWS resources, we should run:

```
```bash
cdk destroy
````
```

For disk cleanup, I can remove caches, virtual environments, and old project folders using commands like:

```
```bash
find ~/.cache -type f -delete
rm -rf .env venv node_modules
pip cache purge
npm cache clean --force
````
```

All the steps of cleaning can be executed using the **destroy.sh** script. I initially ran them step by step, and then combined them into a single batch file. I did not destroy cdk and leave it to be checked by you.

Info: Additional Hands-On

I developed more complex web applications using Machine Learning (ML) and Natural Language Processing (NLP), but because of your request, I deployed simple web application.

- ML Application: [Data Analysis GitHub Repo] (<https://github.com/MaryamSaamin/Data-analysis.git>)

Flask Application -Analytical Service Development and I explained everything in BDA Deployment Project Report-Saamin.pdf, and I run this flask application in **Azure Web App**.

- NLP Application: [NLP GitHub Repo] (<https://github.com/MaryamSaamin/NLP.git>)

